

DEVOPS MASTER COURSE

Volume 13B

Backstage Developer Portal Engineering

Architecture | Catalog | Templates | TechDocs | Plugins | Permissions | Reliability

Chapters 11-20 | 2026 Edition

Beginner to Enterprise Mastery

About This Part

Part B develops production depth in Backstage: catalog modeling, ingestion, templates, TechDocs, plugins, permissions, integrations, upgrades and reliable operations.

2026 curriculum note: Backstage remains an open-source framework for building developer portals. The course treats its catalog, templates and TechDocs as components of a wider platform product rather than the platform by themselves.

Volume 13 Roadmap

Part	Coverage	Status
13A	Platform product foundations and IDP architecture	Course sequence
13B	Backstage developer portal engineering	Current part
13C	Self-service infrastructure and control planes	Course sequence
13D	Golden paths and developer experience	Course sequence
13E	Enterprise platform operations, governance and scale	Course sequence
13F	Projects, interviews and final platform capstone	Course sequence

Scope, Safety and Professional Practice

- Use disposable sandboxes, local simulations or approved training subscriptions. Configure budgets, quotas and expiry before creating resources.
- Never use production identities, customer data, employer account structures, internal addresses or confidential contracts in course artifacts.
- Open-source projects, cloud services, AI models, pricing, limits and certification objectives change. Verify official documentation before production use or exam registration.
- Run failure, security, red-team and recovery tests only in approved nonproduction environments with explicit stop conditions.
- Destroy resources, revoke credentials, remove local state and sanitize every portfolio artifact.

Learning Outcomes

- Explain and design the capabilities covered by Backstage Developer Portal Engineering.
- Implement or simulate production workflows through version-controlled automation.
- Test identity, policy, security, reliability, cost and lifecycle behavior.
- Troubleshoot complex failures using a structured evidence-first method.
- Produce sanitized architecture, laboratory, runbook and interview portfolio artifacts.

Table of Contents

- Chapter 11 Backstage Architecture and Deployment
- Chapter 12 Software Catalog Data Model
- Chapter 13 Catalog Ingestion, Discovery and Quality
- Chapter 14 Software Templates and Scaffolded Workflows
- Chapter 15 TechDocs and Documentation as Code
- Chapter 16 Backstage Plugins, Search and Extensibility
- Chapter 17 Authentication, Permissions and Multi-Tenancy
- Chapter 18 Portal Integrations and Operational Views
- Chapter 19 Backstage Reliability, Security and Upgrades
- Chapter 20 Enterprise Developer Portal Capstone
- Backstage Engineering and Operations Reference
- Completion Checklist
- Additional Portfolio Projects
- Practice Plan
- Glossary
- Official Documentation Index

Chapter 11: Backstage Architecture and Deployment

Understand frontend, backend, plugins, database, authentication and deployment boundaries.

Why this chapter matters: Backstage Architecture and Deployment matters because enterprise outcomes depend on more than a successful tool installation. The team must make understand frontend, backend, plugins, database, authentication and deployment boundaries. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of backstage architecture and deployment.
- Design the capability using Backstage, Node.js, PostgreSQL, Kubernetes with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: deploy a local backstage instance with production-style configuration.

Core Concepts

Understand frontend and problem framing

Understand frontend and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For backstage architecture and deployment, record assumptions and the evidence that would disprove them.

backend architecture and operating boundaries

The architecture for backend architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

plugins automation, policy and evidence

Automation for plugins automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

database failure handling and continuous improvement

Continuous improvement for database failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Backstage	Implement the core capability and define its service contract.
Automation and integration	Node.js	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	PostgreSQL	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubernetes	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Backstage	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	Node.js	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: backstage_architecture_and_deploy
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 11

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 11: Deploy a local Backstage instance with production-style configuration

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for backstage architecture and deployment.
2. Create the smallest safe implementation or detailed simulation using Backstage, Node.js, PostgreSQL, Kubernetes.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Backstage Architecture and Deployment?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Backstage Architecture and Deployment is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 12: Software Catalog Data Model

Model components, APIs, resources, systems, domains, groups and users with durable ownership.

Why this chapter matters: Software Catalog Data Model matters because enterprise outcomes depend on more than a successful tool installation. The team must make model components, apis, resources, systems, domains, groups and users with durable ownership. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of software catalog data model.
- Design the capability using catalog-info.yaml, entity relations, processors with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: create and validate a multi-system catalog model.

Core Concepts

Model components and problem framing

Model components and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For software catalog data model, record assumptions and the evidence that would disprove them.

APIs architecture and operating boundaries

The architecture for apis architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

resources automation, policy and evidence

Automation for resources automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

systems failure handling and continuous improvement

Continuous improvement for systems failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	catalog-info.yaml	Implement the core capability and define its service contract.
Automation and integration	entity relations	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	processors	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	catalog-info.yaml	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	entity relations	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	processors	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: software_catalog_data_model
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 12

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 12: Create and validate a multi-system catalog model

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for software catalog data model.
2. Create the smallest safe implementation or detailed simulation using catalog-info.yaml, entity relations, processors.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Software Catalog Data Model?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Software Catalog Data Model is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 13: Catalog Ingestion, Discovery and Quality

Automate discovery while preventing stale, duplicated or unowned entities.

Why this chapter matters: Catalog Ingestion, Discovery and Quality matters because enterprise outcomes depend on more than a successful tool installation. The team must make automate discovery while preventing stale, duplicated or unowned entities. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of catalog ingestion, discovery and quality.
- Design the capability using GitHub/GitLab discovery, catalog processors, scorecards with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build an ingestion pipeline with metadata quality checks.

Core Concepts

Automate discovery while preventing stale and problem framing

Automate discovery while preventing stale and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For catalog ingestion, discovery and quality, record assumptions and the evidence that would disprove them.

duplicated or unowned entities. architecture and operating boundaries

The architecture for duplicated or unowned entities. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Catalog Ingestion, Discovery and Quality automation, policy and evidence

Automation for catalog ingestion, discovery and quality automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Catalog Ingestion, Discovery and Quality failure handling and continuous improvement

Continuous improvement for catalog ingestion, discovery and quality failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	GitHub/GitLab discovery	Implement the core capability and define its service contract.
Automation and integration	catalog processors	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	scorecards	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	GitHub/GitLab discovery	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	catalog processors	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	scorecards	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: catalog_ingestion_discovery_and_
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 13

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 13: Build an ingestion pipeline with metadata quality checks

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for catalog ingestion, discovery and quality.
2. Create the smallest safe implementation or detailed simulation using GitHub/GitLab discovery, catalog processors, scorecards.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Catalog Ingestion, Discovery and Quality?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Catalog Ingestion, Discovery and Quality is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 14: Software Templates and Scaffolder Workflows

Create safe self-service templates that generate repositories, pipelines and environments.

Why this chapter matters: Software Templates and Scaffolder Workflows matters because enterprise outcomes depend on more than a successful tool installation. The team must make create safe self-service templates that generate repositories, pipelines and environments. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of software templates and scaffolder workflows.
- Design the capability using Backstage Software Templates, actions, secrets with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build a service template with validation and rollback.

Core Concepts

Create safe self-service templates that generate repositories and problem framing

Create safe self-service templates that generate repositories and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For software templates and scaffolder workflows, record assumptions and the evidence that would disprove them.

pipelines and environments. architecture and operating boundaries

The architecture for pipelines and environments. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Software Templates and Scaffolder Workflows automation, policy and evidence

Automation for software templates and scaffolder workflows automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Software Templates and Scaffolder Workflows failure handling and continuous improvement

Continuous improvement for software templates and scaffolder workflows failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Backstage Software Templates	Implement the core capability and define its service contract.
Automation and integration	actions	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	secrets	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Backstage Software Templates	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	actions	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	secrets	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: software_templates_and_scaffolds
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 14

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 14: Build a service template with validation and rollback

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for software templates and scaffolder workflows.
2. Create the smallest safe implementation or detailed simulation using Backstage Software Templates, actions, secrets.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Software Templates and Scaffolder Workflows?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Software Templates and Scaffolder Workflows is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 15: TechDocs and Documentation as Code

Publish discoverable, versioned documentation close to the software it describes.

Why this chapter matters: TechDocs and Documentation as Code matters because enterprise outcomes depend on more than a successful tool installation. The team must make publish discoverable, versioned documentation close to the software it describes. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of techdocs and documentation as code.
- Design the capability using TechDocs, MkDocs, object storage with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: publish techdocs for a sample service and validate ownership.

Core Concepts

Publish discoverable and problem framing

Publish discoverable and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For techdocs and documentation as code, record assumptions and the evidence that would disprove them.

versioned documentation close to the software it describes. architecture and operating boundaries

The architecture for versioned documentation close to the software it describes. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

TechDocs and Documentation as Code automation, policy and evidence

Automation for techdocs and documentation as code automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

TechDocs and Documentation as Code failure handling and continuous improvement

Continuous improvement for techdocs and documentation as code failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	TechDocs	Implement the core capability and define its service contract.
Automation and integration	MkDocs	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	object storage	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	TechDocs	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	MkDocs	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	object storage	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: techdocs_and_documentation_as_co
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 15

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 15: Publish TechDocs for a sample service and validate ownership

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for techdocs and documentation as code.
2. Create the smallest safe implementation or detailed simulation using TechDocs, MkDocs, object storage.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design TechDocs and Documentation as Code?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- TechDocs and Documentation as Code is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 16: Backstage Plugins, Search and Extensibility

Extend the portal without creating an unmaintainable plugin estate.

Why this chapter matters: Backstage Plugins, Search and Extensibility matters because enterprise outcomes depend on more than a successful tool installation. The team must make extend the portal without creating an unmaintainable plugin estate. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of backstage plugins, search and extensibility.
- Design the capability using Backstage plugins, search, APIs, analytics with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: create a small plugin and integrate it with catalog entities.

Core Concepts

Extend the portal without creating an unmaintainable plugin estate. and problem framing

Extend the portal without creating an unmaintainable plugin estate. and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For backstage plugins, search and extensibility, record assumptions and the evidence that would disprove them.

Backstage Plugins, Search and Extensibility architecture and operating boundaries

The architecture for backstage plugins, search and extensibility architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Backstage Plugins, Search and Extensibility automation, policy and evidence

Automation for backstage plugins, search and extensibility automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Backstage Plugins, Search and Extensibility failure handling and continuous improvement

Continuous improvement for backstage plugins, search and extensibility failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Backstage plugins	Implement the core capability and define its service contract.
Automation and integration	search	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	APIs	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	analytics	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	Backstage plugins	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	search	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: backstage_plugins_search_and_ext
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 16

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 16: Create a small plugin and integrate it with catalog entities

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for backstage plugins, search and extensibility.
2. Create the smallest safe implementation or detailed simulation using Backstage plugins, search, APIs, analytics.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Backstage Plugins, Search and Extensibility?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Backstage Plugins, Search and Extensibility is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 17: Authentication, Permissions and Multi-Tenancy

Federate identity and enforce least-privilege portal actions and data visibility.

Why this chapter matters: Authentication, Permissions and Multi-Tenancy matters because enterprise outcomes depend on more than a successful tool installation. The team must make federate identity and enforce least-privilege portal actions and data visibility. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of authentication, permissions and multi-tenancy.
- Design the capability using OIDC, Backstage permissions framework, RBAC with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: implement role-aware access for platform and product teams.

Core Concepts

Federate identity and enforce least-privilege portal actions and data visibility. and problem framing

Federate identity and enforce least-privilege portal actions and data visibility. and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For authentication, permissions and multi-tenancy, record assumptions and the evidence that would disprove them.

Authentication, Permissions and Multi-Tenancy architecture and operating boundaries

The architecture for authentication, permissions and multi-tenancy architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Authentication, Permissions and Multi-Tenancy automation, policy and evidence

Automation for authentication, permissions and multi-tenancy automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Authentication, Permissions and Multi-Tenancy failure handling and continuous improvement

Continuous improvement for authentication, permissions and multi-tenancy failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	OIDC	Implement the core capability and define its service contract.
Automation and integration	Backstage permissions framework	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	RBAC	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	OIDC	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	Backstage permissions framework	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	RBAC	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: authentication_permissions_and_m
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 17

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 17: Implement role-aware access for platform and product teams

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for authentication, permissions and multi-tenancy.
2. Create the smallest safe implementation or detailed simulation using OIDC, Backstage permissions framework, RBAC.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Authentication, Permissions and Multi-Tenancy?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Authentication, Permissions and Multi-Tenancy is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 18: Portal Integrations and Operational Views

Integrate CI/CD, observability, cloud, Kubernetes and incident tooling around service context.

Why this chapter matters: Portal Integrations and Operational Views matters because enterprise outcomes depend on more than a successful tool installation. The team must make integrate ci/cd, observability, cloud, kubernetes and incident tooling around service context. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of portal integrations and operational views.
- Design the capability using GitHub Actions, Argo CD, Grafana, Kubernetes plugins with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build a service page that links delivery and runtime evidence.

Core Concepts

Integrate CI/CD and problem framing

Integrate CI/CD and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For portal integrations and operational views, record assumptions and the evidence that would disprove them.

observability architecture and operating boundaries

The architecture for observability architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

cloud automation, policy and evidence

Automation for cloud automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Kubernetes and incident tooling around service context. failure handling and continuous improvement

Continuous improvement for kubernetes and incident tooling around service context. failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	GitHub Actions	Implement the core capability and define its service contract.
Automation and integration	Argo CD	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Grafana	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubernetes plugins	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	GitHub Actions	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	Argo CD	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: portal_integrations_and_operatio
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 18

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 18: Build a service page that links delivery and runtime evidence

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for portal integrations and operational views.
2. Create the smallest safe implementation or detailed simulation using GitHub Actions, Argo CD, Grafana, Kubernetes plugins.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Portal Integrations and Operational Views?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Portal Integrations and Operational Views is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 19: Backstage Reliability, Security and Upgrades

Operate the portal with backup, scaling, dependency, vulnerability and upgrade controls.

Why this chapter matters: Backstage Reliability, Security and Upgrades matters because enterprise outcomes depend on more than a successful tool installation. The team must make operate the portal with backup, scaling, dependency, vulnerability and upgrade controls. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of backstage reliability, security and upgrades.
- Design the capability using Kubernetes, PostgreSQL HA, SLOs, dependency scanning with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: run an upgrade and database recovery exercise.

Core Concepts

Operate the portal with backup and problem framing

Operate the portal with backup and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For backstage reliability, security and upgrades, record assumptions and the evidence that would disprove them.

scaling architecture and operating boundaries

The architecture for scaling architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

dependency automation, policy and evidence

Automation for dependency automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

vulnerability and upgrade controls. failure handling and continuous improvement

Continuous improvement for vulnerability and upgrade controls. failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubernetes	Implement the core capability and define its service contract.
Automation and integration	PostgreSQL HA	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	SLOs	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	dependency scanning	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Kubernetes	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	PostgreSQL HA	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: backstage_reliability_security_a
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 19

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 19: Run an upgrade and database recovery exercise

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for backstage reliability, security and upgrades.
2. Create the smallest safe implementation or detailed simulation using Kubernetes, PostgreSQL HA, SLOs, dependency scanning.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Backstage Reliability, Security and Upgrades?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Backstage Reliability, Security and Upgrades is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 20: Enterprise Developer Portal Capstone

Deliver a governed portal with catalog, templates, docs, permissions and operational integrations.

Why this chapter matters: Enterprise Developer Portal Capstone matters because enterprise outcomes depend on more than a successful tool installation. The team must make deliver a governed portal with catalog, templates, docs, permissions and operational integrations. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of enterprise developer portal capstone.
- Design the capability using Backstage, Kubernetes, GitOps, OpenTelemetry with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: deliver a production-readiness review for an enterprise portal.

Core Concepts

Deliver a governed portal with catalog and problem framing

Deliver a governed portal with catalog and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For enterprise developer portal capstone, record assumptions and the evidence that would disprove them.

templates architecture and operating boundaries

The architecture for templates architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

docs automation, policy and evidence

Automation for docs automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

permissions and operational integrations. failure handling and continuous improvement

Continuous improvement for permissions and operational integrations. failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Backstage	Implement the core capability and define its service contract.
Automation and integration	Kubernetes	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	GitOps	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	OpenTelemetry	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	Backstage	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	Kubernetes	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: enterprise_developer_portal_caps
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 13B
chapter: 20

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 20: Deliver a production-readiness review for an enterprise portal

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for enterprise developer portal capstone.
2. Create the smallest safe implementation or detailed simulation using Backstage, Kubernetes, GitOps, OpenTelemetry.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Enterprise Developer Portal Capstone?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Enterprise Developer Portal Capstone is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Backstage Engineering and Operations Reference

Use this reference to review product outcomes, architecture boundaries, operating controls and evidence required for production approval.

Domain	Required outcome	Validation evidence
Product and users	Named consumers, outcome hypothesis, roadmap and lifecycle owner	Research summary, service contract and adoption data
Architecture	Control, execution, data and evidence boundaries with bounded blast radius	Diagrams, ADRs, dependency and failure tests
Identity and security	Least privilege, keyless workloads, protected administration and auditable exceptions	Effective-access, denied-action and audit evidence
Delivery and policy	Versioned automation, policy, promotion, rollback and provenance	Pipeline, policy and artifact verification results
Reliability and operations	SLOs, telemetry, capacity, incident response, recovery and deprecation	Dashboards, game-day timeline and restore evidence
Data and AI governance	Ownership, classification, quality, privacy, lineage and responsible use	Catalog, tests, risk and approval records
Economics and value	Allocated cost, unit economics, investment rationale and optimization safety	Cost model, outcome metrics and decision log
People and governance	Clear decisions, support, skills, change plan and feedback loops	RACI, operating cadence, survey and action backlog

Completion Checklist

- [] I can explain every chapter in Volume 13B using architecture, operations and business language.
- [] I can implement or simulate at least one complete service with identity, policy, telemetry and lifecycle controls.
- [] I can demonstrate expected, denied and failure behavior with time-stamped evidence.
- [] I can discuss security, reliability, cost, adoption and organizational trade-offs.
- [] I can produce a sanitized portfolio artifact and defend the design in an interview or review board.

Additional Portfolio Projects

1. Service catalog and ownership quality automation.
2. Keyless CI/CD and workload-identity reference implementation.
3. Policy promotion pipeline with expiring exceptions.
4. Golden path or pipeline template with user-journey telemetry.
5. SLO dashboard and bounded failure game day.
6. Cost allocation and unit-economics review.
7. Security threat model and evidence-based remediation plan.
8. Backup, restore, failover and failback exercise.
9. Migration or adoption factory for three fictional teams.
10. Executive architecture or investment memo with measured outcomes.

Practice Plan

Days	Focus	Required output
1-5	Concepts, users and current-state discovery	Research notes, value stream and capability map
6-10	Architecture and service contracts	Diagrams, API/schema and ADRs
11-16	Implementation and automation	Working sandbox or detailed simulation
17-21	Security, policy and governance	Threat model, policy tests and exceptions
22-26	Reliability, telemetry and recovery	SLOs, dashboards and failure drill
27-31	Cost, capacity and lifecycle	Unit cost, quota and deprecation plan
32-36	Adoption, people and operating model	RACI, support and change plan
37-41	Troubleshooting and interview practice	Timed investigation and interview notes
42-45	Capstone and portfolio	Final defense, cleanup and sanitized evidence pack

Glossary

Term	Definition
Control plane	The APIs and controllers that accept desired state and coordinate execution.
Developer portal	A user interface that aggregates catalog, documentation, workflows and operational context.
Internal developer platform	The integrated products, APIs, automation and processes that support software delivery.
Golden path	An opinionated, supported default journey for a common engineering task.
Data product	An owned data asset or interface with consumers, quality, documentation and lifecycle.
MLOps	Practices and systems for reproducible, governed delivery and operation of machine-learning models.
LLMOps	Lifecycle, evaluation, deployment, observability, safety and cost practices for LLM applications.
Sociotechnical system	A system whose outcomes emerge from interactions among people, process, incentives and technology.
Unit economics	Cost and value measured per meaningful business or technical output.
Error budget	The allowed unreliability implied by a service-level objective.
Policy as code	Machine-testable policy stored, reviewed and promoted through version control.
Evidence pack	Sanitized architecture, configuration, test, telemetry and decision records used to prove readiness.

Official Documentation Index

The following primary or official sources were consulted for this 2026 edition. Verify current versions, service behavior, prices, limits and certification objectives before use.

- [Backstage - What is Backstage?](#)
- [Backstage Software Catalog](#)
- [Backstage Software Templates](#)
- [Backstage TechDocs](#)

- [Crossplane Documentation](#)
- [Crossplane Compositions](#)
- [CNCF Platforms White Paper](#)
- [DORA Platform Engineering Capability](#)
- [Kubernetes Documentation](#)
- [OpenGitOps Principles](#)

Volume 13B completion: preserve your architecture, code or simulation, test evidence, recovery timeline, cost/value review and sanitized final report as one portfolio package.