

DEVOPS MASTER COURSE

Volume 14B

MLOps Lifecycle Engineering

*Tracking | Datasets | Features | Pipelines | Registry | CI/CT/CD | Evaluation |
Monitoring*

Chapters 11-20 | 2026 Edition

Beginner to Enterprise Mastery

About This Part

Part B creates a governed model lifecycle from problem framing and reproducible experiments through registry, delivery, evaluation and monitoring.

2026 curriculum note: Current MLflow documentation covers both traditional ML and LLM/agent workflows. This part focuses on the traditional model lifecycle while retaining lineage and evaluation practices that carry forward to LLMOps.

Volume 14 Roadmap

Part	Coverage	Status
14A	DataOps and modern data-platform foundations	Course sequence
14B	MLOps lifecycle, registry and monitoring	Current part
14C	Kubeflow, Ray and distributed AI infrastructure	Course sequence
14D	Model serving, RAG, agents and LLMOps	Course sequence
14E	AI infrastructure operations, security, risk and FinOps	Course sequence
14F	Projects, interviews and final AI platform capstone	Course sequence

Scope, Safety and Professional Practice

- Use disposable sandboxes, local simulations or approved training subscriptions. Configure budgets, quotas and expiry before creating resources.
- Never use production identities, customer data, employer account structures, internal addresses or confidential contracts in course artifacts.
- Open-source projects, cloud services, AI models, pricing, limits and certification objectives change. Verify official documentation before production use or exam registration.
- Run failure, security, red-team and recovery tests only in approved nonproduction environments with explicit stop conditions.
- Destroy resources, revoke credentials, remove local state and sanitize every portfolio artifact.

Learning Outcomes

- Explain and design the capabilities covered by MLOps Lifecycle Engineering.
- Implement or simulate production workflows through version-controlled automation.
- Test identity, policy, security, reliability, cost and lifecycle behavior.
- Troubleshoot complex failures using a structured evidence-first method.
- Produce sanitized architecture, laboratory, runbook and interview portfolio artifacts.

Table of Contents

- Chapter 11 ML Product Framing and Success Criteria
- Chapter 12 Experiment Tracking with MLflow
- Chapter 13 Dataset Versioning and Lineage for ML
- Chapter 14 Feature Engineering and Feature Stores
- Chapter 15 ML Pipeline Orchestration
- Chapter 16 Model Registry and Promotion
- Chapter 17 Continuous Integration, Training and Delivery
- Chapter 18 Model Evaluation and Validation
- Chapter 19 Model Monitoring, Drift and Feedback
- Chapter 20 Production MLOps Capstone
- MLOps Lifecycle and Governance Reference
- Completion Checklist
- Additional Portfolio Projects
- Practice Plan
- Glossary
- Official Documentation Index

Chapter 11: ML Product Framing and Success Criteria

Translate business decisions into measurable ML objectives, baselines and safe fallback behavior.

Why this chapter matters: ML Product Framing and Success Criteria matters because enterprise outcomes depend on more than a successful tool installation. The team must make translate business decisions into measurable ml objectives, baselines and safe fallback behavior. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of ml product framing and success criteria.
- Design the capability using Problem canvases, baselines, cost-of-error with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: create an ml product brief and evaluation plan.

Core Concepts

Translate business decisions into measurable ML objectives and problem framing

Translate business decisions into measurable ML objectives and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For ml product framing and success criteria, record assumptions and the evidence that would disprove them.

baselines and safe fallback behavior. architecture and operating boundaries

The architecture for baselines and safe fallback behavior. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

ML Product Framing and Success Criteria automation, policy and evidence

Automation for ml product framing and success criteria automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

ML Product Framing and Success Criteria failure handling and continuous improvement

Continuous improvement for ml product framing and success criteria failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Problem canvases	Implement the core capability and define its service contract.
Automation and integration	baselines	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	cost-of-error	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Problem canvases	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	baselines	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	cost-of-error	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: ml_product_framing_and_success_c
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 11

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 11: Create an ML product brief and evaluation plan

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for ml product framing and success criteria.
2. Create the smallest safe implementation or detailed simulation using Problem canvases, baselines, cost-of-error.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design ML Product Framing and Success Criteria?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- ML Product Framing and Success Criteria is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 12: Experiment Tracking with MLflow

Capture code, parameters, metrics, artifacts and reproducibility evidence.

Why this chapter matters: Experiment Tracking with MLflow matters because enterprise outcomes depend on more than a successful tool installation. The team must make capture code, parameters, metrics, artifacts and reproducibility evidence. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of experiment tracking with mlflow.
- Design the capability using MLflow Tracking, artifact stores, databases with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: track and compare reproducible experiments.

Core Concepts

Capture code and problem framing

Capture code and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For experiment tracking with mlflow, record assumptions and the evidence that would disprove them.

parameters architecture and operating boundaries

The architecture for parameters architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

metrics automation, policy and evidence

Automation for metrics automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

artifacts and reproducibility evidence. failure handling and continuous improvement

Continuous improvement for artifacts and reproducibility evidence. failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	MLflow Tracking	Implement the core capability and define its service contract.
Automation and integration	artifact stores	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	databases	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	MLflow Tracking	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	artifact stores	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	databases	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: experiment_tracking_with_mlflow
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 12

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 12: Track and compare reproducible experiments

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for experiment tracking with mlflow.
2. Create the smallest safe implementation or detailed simulation using MLflow Tracking, artifact stores, databases.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Experiment Tracking with MLflow?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Experiment Tracking with MLflow is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 13: Dataset Versioning and Lineage for ML

Connect training, validation and evaluation data to model outcomes and governance.

Why this chapter matters: Dataset Versioning and Lineage for ML matters because enterprise outcomes depend on more than a successful tool installation. The team must make connect training, validation and evaluation data to model outcomes and governance. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of dataset versioning and lineage for ml.
- Design the capability using MLflow datasets, DVC, lakehouse snapshots with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build a dataset lineage and versioning workflow.

Core Concepts

Connect training and problem framing

Connect training and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For dataset versioning and lineage for ml, record assumptions and the evidence that would disprove them.

validation and evaluation data to model outcomes and governance. architecture and operating boundaries

The architecture for validation and evaluation data to model outcomes and governance. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Dataset Versioning and Lineage for ML automation, policy and evidence

Automation for dataset versioning and lineage for ml automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Dataset Versioning and Lineage for ML failure handling and continuous improvement

Continuous improvement for dataset versioning and lineage for ml failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	MLflow datasets	Implement the core capability and define its service contract.
Automation and integration	DVC	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	lakehouse snapshots	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	MLflow datasets	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	DVC	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	lakehouse snapshots	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: dataset_versioning_and_lineage_f
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 13

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 13: Build a dataset lineage and versioning workflow

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for dataset versioning and lineage for ml.
2. Create the smallest safe implementation or detailed simulation using MLflow datasets, DVC, lakehouse snapshots.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Dataset Versioning and Lineage for ML?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Dataset Versioning and Lineage for ML is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 14: Feature Engineering and Feature Stores

Prevent training-serving skew and manage reusable offline and online features.

Why this chapter matters: Feature Engineering and Feature Stores matters because enterprise outcomes depend on more than a successful tool installation. The team must make prevent training-serving skew and manage reusable offline and online features. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of feature engineering and feature stores.
- Design the capability using Feast, feature pipelines, point-in-time joins with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: create and validate a feature service.

Core Concepts

Prevent training-serving skew and manage reusable offline and online features. and problem framing

Prevent training-serving skew and manage reusable offline and online features. and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For feature engineering and feature stores, record assumptions and the evidence that would disprove them.

Feature Engineering and Feature Stores architecture and operating boundaries

The architecture for feature engineering and feature stores architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Feature Engineering and Feature Stores automation, policy and evidence

Automation for feature engineering and feature stores automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Feature Engineering and Feature Stores failure handling and continuous improvement

Continuous improvement for feature engineering and feature stores failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Feast	Implement the core capability and define its service contract.
Automation and integration	feature pipelines	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	point-in-time joins	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Feast	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	feature pipelines	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	point-in-time joins	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: feature_engineering_and_feature_
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 14

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 14: Create and validate a feature service

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for feature engineering and feature stores.
2. Create the smallest safe implementation or detailed simulation using Feast, feature pipelines, point-in-time joins.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Feature Engineering and Feature Stores?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Feature Engineering and Feature Stores is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 15: ML Pipeline Orchestration

Package preprocessing, training, evaluation and registration as repeatable workflows.

Why this chapter matters: ML Pipeline Orchestration matters because enterprise outcomes depend on more than a successful tool installation. The team must make package preprocessing, training, evaluation and registration as repeatable workflows. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of ml pipeline orchestration.
- Design the capability using Kubeflow Pipelines, Airflow, Argo Workflows with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build an end-to-end training pipeline.

Core Concepts

Package preprocessing and problem framing

Package preprocessing and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For ml pipeline orchestration, record assumptions and the evidence that would disprove them.

training architecture and operating boundaries

The architecture for training architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

evaluation and registration as repeatable workflows. automation, policy and evidence

Automation for evaluation and registration as repeatable workflows. automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

ML Pipeline Orchestration failure handling and continuous improvement

Continuous improvement for ml pipeline orchestration failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubeflow Pipelines	Implement the core capability and define its service contract.
Automation and integration	Airflow	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Argo Workflows	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubeflow Pipelines	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Airflow	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	Argo Workflows	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: ml_pipeline_orchestration
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 15

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 15: Build an end-to-end training pipeline

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for ml pipeline orchestration.
2. Create the smallest safe implementation or detailed simulation using Kubeflow Pipelines, Airflow, Argo Workflows.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design ML Pipeline Orchestration?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- ML Pipeline Orchestration is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 16: Model Registry and Promotion

Manage lineage, versions, aliases, approvals, signatures and lifecycle.

Why this chapter matters: Model Registry and Promotion matters because enterprise outcomes depend on more than a successful tool installation. The team must manage lineage, versions, aliases, approvals, signatures and lifecycle. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of model registry and promotion.
- Design the capability using MLflow Model Registry, metadata, policy with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: implement a staged model promotion workflow.

Core Concepts

Manage lineage and problem framing

Manage lineage and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For model registry and promotion, record assumptions and the evidence that would disprove them.

versions architecture and operating boundaries

The architecture for versions architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

aliases automation, policy and evidence

Automation for aliases automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

approvals failure handling and continuous improvement

Continuous improvement for approvals failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	MLflow Model Registry	Implement the core capability and define its service contract.
Automation and integration	metadata	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	policy	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	MLflow Model Registry	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	metadata	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	policy	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: model_registry_and_promotion
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 16

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 16: Implement a staged model promotion workflow

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for model registry and promotion.
2. Create the smallest safe implementation or detailed simulation using MLflow Model Registry, metadata, policy.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Model Registry and Promotion?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Model Registry and Promotion is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 17: Continuous Integration, Training and Delivery

Automate code, data, pipeline and model validation without blindly retraining.

Why this chapter matters: Continuous Integration, Training and Delivery matters because enterprise outcomes depend on more than a successful tool installation. The team must make automate code, data, pipeline and model validation without blindly retraining. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of continuous integration, training and delivery.
- Design the capability using CI/CD, CT, IaC, GitOps with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build a model delivery pipeline with approval gates.

Core Concepts

Automate code and problem framing

Automate code and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For continuous integration, training and delivery, record assumptions and the evidence that would disprove them.

data architecture and operating boundaries

The architecture for data architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

pipeline and model validation without blindly retraining. automation, policy and evidence

Automation for pipeline and model validation without blindly retraining. automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Continuous Integration, Training and Delivery failure handling and continuous improvement

Continuous improvement for continuous integration, training and delivery failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	CI/CD	Implement the core capability and define its service contract.
Automation and integration	CT	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	IaC	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	GitOps	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	CI/CD	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	CT	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: continuous_integration_training_
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 17

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 17: Build a model delivery pipeline with approval gates

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for continuous integration, training and delivery.
2. Create the smallest safe implementation or detailed simulation using CI/CD, CT, IaC, GitOps.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Continuous Integration, Training and Delivery?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Continuous Integration, Training and Delivery is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 18: Model Evaluation and Validation

Test predictive quality, robustness, fairness, latency and business impact.

Why this chapter matters: Model Evaluation and Validation matters because enterprise outcomes depend on more than a successful tool installation. The team must make test predictive quality, robustness, fairness, latency and business impact. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of model evaluation and validation.
- Design the capability using MLflow evaluation, custom metrics, test sets with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: create a model evaluation report and release gate.

Core Concepts

Test predictive quality and problem framing

Test predictive quality and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For model evaluation and validation, record assumptions and the evidence that would disprove them.

robustness architecture and operating boundaries

The architecture for robustness architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

fairness automation, policy and evidence

Automation for fairness automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

latency and business impact. failure handling and continuous improvement

Continuous improvement for latency and business impact. failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	MLflow evaluation	Implement the core capability and define its service contract.
Automation and integration	custom metrics	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	test sets	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	MLflow evaluation	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	custom metrics	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	test sets	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: model_evaluation_and_validation
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 18

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 18: Create a model evaluation report and release gate

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for model evaluation and validation.
2. Create the smallest safe implementation or detailed simulation using MLflow evaluation, custom metrics, test sets.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Model Evaluation and Validation?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Model Evaluation and Validation is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 19: Model Monitoring, Drift and Feedback

Detect data, concept, performance and operational degradation with actionable response.

Why this chapter matters: Model Monitoring, Drift and Feedback matters because enterprise outcomes depend on more than a successful tool installation. The team must make detect data, concept, performance and operational degradation with actionable response. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of model monitoring, drift and feedback.
- Design the capability using Evidently, Prometheus, OpenTelemetry with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build a drift dashboard and retraining decision policy.

Core Concepts

Detect data and problem framing

Detect data and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For model monitoring, drift and feedback, record assumptions and the evidence that would disprove them.

concept architecture and operating boundaries

The architecture for concept architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

performance and operational degradation with actionable response. automation, policy and evidence

Automation for performance and operational degradation with actionable response. automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Model Monitoring, Drift and Feedback failure handling and continuous improvement

Continuous improvement for model monitoring, drift and feedback failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Evidently	Implement the core capability and define its service contract.
Automation and integration	Prometheus	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	OpenTelemetry	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Evidently	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Prometheus	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	OpenTelemetry	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: model_monitoring_drift_and_feedb
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 19

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 19: Build a drift dashboard and retraining decision policy

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for model monitoring, drift and feedback.
2. Create the smallest safe implementation or detailed simulation using Evidently, Prometheus, OpenTelemetry.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Model Monitoring, Drift and Feedback?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Model Monitoring, Drift and Feedback is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 20: Production MLOps Capstone

Integrate tracking, datasets, features, pipelines, registry, deployment and monitoring.

Why this chapter matters: Production MLOps Capstone matters because enterprise outcomes depend on more than a successful tool installation. The team must make integrate tracking, datasets, features, pipelines, registry, deployment and monitoring. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of production mlops capstone.
- Design the capability using MLflow, Kubeflow, Feast, KServe with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: deliver a governed mlops lifecycle.

Core Concepts

Integrate tracking and problem framing

Integrate tracking and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For production mlops capstone, record assumptions and the evidence that would disprove them.

datasets architecture and operating boundaries

The architecture for datasets architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

features automation, policy and evidence

Automation for features automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

pipelines failure handling and continuous improvement

Continuous improvement for pipelines failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	MLflow	Implement the core capability and define its service contract.
Automation and integration	Kubeflow	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Feast	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	KServe	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	MLflow	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	Kubeflow	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: production_mlops_capstone
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14B
chapter: 20

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 20: Deliver a governed MLOps lifecycle

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for production mlops capstone.
2. Create the smallest safe implementation or detailed simulation using MLflow, Kubeflow, Feast, KServe.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Production MLOps Capstone?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Production MLOps Capstone is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

MLOps Lifecycle and Governance Reference

Use this reference to review product outcomes, architecture boundaries, operating controls and evidence required for production approval.

Domain	Required outcome	Validation evidence
Product and users	Named consumers, outcome hypothesis, roadmap and lifecycle owner	Research summary, service contract and adoption data
Architecture	Control, execution, data and evidence boundaries with bounded blast radius	Diagrams, ADRs, dependency and failure tests
Identity and security	Least privilege, keyless workloads, protected administration and auditable exceptions	Effective-access, denied-action and audit evidence
Delivery and policy	Versioned automation, policy, promotion, rollback and provenance	Pipeline, policy and artifact verification results
Reliability and operations	SLOs, telemetry, capacity, incident response, recovery and deprecation	Dashboards, game-day timeline and restore evidence
Data and AI governance	Ownership, classification, quality, privacy, lineage and responsible use	Catalog, tests, risk and approval records
Economics and value	Allocated cost, unit economics, investment rationale and optimization safety	Cost model, outcome metrics and decision log
People and governance	Clear decisions, support, skills, change plan and feedback loops	RACI, operating cadence, survey and action backlog

Completion Checklist

- [] I can explain every chapter in Volume 14B using architecture, operations and business language.
- [] I can implement or simulate at least one complete service with identity, policy, telemetry and lifecycle controls.
- [] I can demonstrate expected, denied and failure behavior with time-stamped evidence.
- [] I can discuss security, reliability, cost, adoption and organizational trade-offs.
- [] I can produce a sanitized portfolio artifact and defend the design in an interview or review board.

Additional Portfolio Projects

1. Service catalog and ownership quality automation.
2. Keyless CI/CD and workload-identity reference implementation.
3. Policy promotion pipeline with expiring exceptions.
4. Golden path or pipeline template with user-journey telemetry.
5. SLO dashboard and bounded failure game day.
6. Cost allocation and unit-economics review.
7. Security threat model and evidence-based remediation plan.
8. Backup, restore, failover and failback exercise.
9. Migration or adoption factory for three fictional teams.
10. Executive architecture or investment memo with measured outcomes.

Practice Plan

Days	Focus	Required output
1-5	Concepts, users and current-state discovery	Research notes, value stream and capability map
6-10	Architecture and service contracts	Diagrams, API/schema and ADRs
11-16	Implementation and automation	Working sandbox or detailed simulation
17-21	Security, policy and governance	Threat model, policy tests and exceptions
22-26	Reliability, telemetry and recovery	SLOs, dashboards and failure drill
27-31	Cost, capacity and lifecycle	Unit cost, quota and deprecation plan
32-36	Adoption, people and operating model	RACI, support and change plan
37-41	Troubleshooting and interview practice	Timed investigation and interview notes
42-45	Capstone and portfolio	Final defense, cleanup and sanitized evidence pack

Glossary

Term	Definition
Control plane	The APIs and controllers that accept desired state and coordinate execution.
Developer portal	A user interface that aggregates catalog, documentation, workflows and operational context.
Internal developer platform	The integrated products, APIs, automation and processes that support software delivery.
Golden path	An opinionated, supported default journey for a common engineering task.
Data product	An owned data asset or interface with consumers, quality, documentation and lifecycle.
MLOps	Practices and systems for reproducible, governed delivery and operation of machine-learning models.
LLMOps	Lifecycle, evaluation, deployment, observability, safety and cost practices for LLM applications.
Sociotechnical system	A system whose outcomes emerge from interactions among people, process, incentives and technology.
Unit economics	Cost and value measured per meaningful business or technical output.
Error budget	The allowed unreliability implied by a service-level objective.
Policy as code	Machine-testable policy stored, reviewed and promoted through version control.
Evidence pack	Sanitized architecture, configuration, test, telemetry and decision records used to prove readiness.

Official Documentation Index

The following primary or official sources were consulted for this 2026 edition. Verify current versions, service behavior, prices, limits and certification objectives before use.

- [MLflow Documentation](#)
- [MLflow Tracking](#)
- [MLflow Model Registry](#)
- [Kubeflow Introduction](#)

- [Kubeflow Pipelines](#)
- [Kubeflow Trainer](#)
- [KServe Documentation](#)
- [Ray Documentation](#)
- [Ray Serve](#)
- [NIST AI Risk Management Framework](#)

Volume 14B completion: preserve your architecture, code or simulation, test evidence, recovery timeline, cost/value review and sanitized final report as one portfolio package.