

DEVOPS MASTER COURSE

Volume 14C

Kubeflow, Ray and Distributed AI Infrastructure

Notebooks | Pipelines | Training | Katib | GPUs | Ray Data | Ray Train | KubeRay

Chapters 21-30 | 2026 Edition

Beginner to Enterprise Mastery

About This Part

Part C engineers Kubernetes-native workspaces, pipelines, distributed training, optimization, accelerator scheduling and Ray-based compute.

2026 curriculum note: Current Kubeflow documentation describes a modular foundation of tools for AI platforms on Kubernetes. The course adopts components independently where that reduces operational complexity.

Volume 14 Roadmap

Part	Coverage	Status
14A	DataOps and modern data-platform foundations	Course sequence
14B	MLOps lifecycle, registry and monitoring	Course sequence
14C	Kubeflow, Ray and distributed AI infrastructure	Current part
14D	Model serving, RAG, agents and LLMOps	Course sequence
14E	AI infrastructure operations, security, risk and FinOps	Course sequence
14F	Projects, interviews and final AI platform capstone	Course sequence

Scope, Safety and Professional Practice

- Use disposable sandboxes, local simulations or approved training subscriptions. Configure budgets, quotas and expiry before creating resources.
- Never use production identities, customer data, employer account structures, internal addresses or confidential contracts in course artifacts.
- Open-source projects, cloud services, AI models, pricing, limits and certification objectives change. Verify official documentation before production use or exam registration.
- Run failure, security, red-team and recovery tests only in approved nonproduction environments with explicit stop conditions.
- Destroy resources, revoke credentials, remove local state and sanitize every portfolio artifact.

Learning Outcomes

- Explain and design the capabilities covered by Kubeflow, Ray and Distributed AI Infrastructure.
- Implement or simulate production workflows through version-controlled automation.
- Test identity, policy, security, reliability, cost and lifecycle behavior.
- Troubleshoot complex failures using a structured evidence-first method.
- Produce sanitized architecture, laboratory, runbook and interview portfolio artifacts.

Table of Contents

- Chapter 21 Kubernetes-Native AI Platform Architecture
- Chapter 22 Kubeflow Notebooks and Workspace Isolation
- Chapter 23 Kubeflow Pipelines Engineering
- Chapter 24 Kubeflow Trainer and Distributed Jobs
- Chapter 25 Katib Hyperparameter Optimization
- Chapter 26 GPU Scheduling, Quotas and Topology
- Chapter 27 Ray Data for AI Workloads
- Chapter 28 Ray Train for Distributed Training
- Chapter 29 KubeRay Operations and Cluster Lifecycle
- Chapter 30 Distributed AI Platform Capstone
- Distributed AI Platform Reference
- Completion Checklist
- Additional Portfolio Projects
- Practice Plan
- Glossary
- Official Documentation Index

Chapter 21: Kubernetes-Native AI Platform Architecture

Separate notebook, pipeline, training, serving, storage, identity and governance concerns.

Why this chapter matters: Kubernetes-Native AI Platform Architecture matters because enterprise outcomes depend on more than a successful tool installation. The team must make separate notebook, pipeline, training, serving, storage, identity and governance concerns. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of kubernetes-native ai platform architecture.
- Design the capability using Kubeflow, Kubernetes, object storage with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: design a modular ai platform on kubernetes.

Core Concepts

Separate notebook and problem framing

Separate notebook and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For kubernetes-native ai platform architecture, record assumptions and the evidence that would disprove them.

pipeline architecture and operating boundaries

The architecture for pipeline architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

training automation, policy and evidence

Automation for training automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

serving failure handling and continuous improvement

Continuous improvement for serving failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubeflow	Implement the core capability and define its service contract.
Automation and integration	Kubernetes	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	object storage	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubeflow	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Kubernetes	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	object storage	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: kubernetes_native_ai_platform_ar
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 21

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 21: Design a modular AI platform on Kubernetes

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for kubernetes-native ai platform architecture.
2. Create the smallest safe implementation or detailed simulation using Kubeflow, Kubernetes, object storage.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Kubernetes-Native AI Platform Architecture?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Kubernetes-Native AI Platform Architecture is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 22: Kubeflow Notebooks and Workspace Isolation

Provide reproducible workspaces with quotas, images, data access and lifecycle controls.

Why this chapter matters: Kubeflow Notebooks and Workspace Isolation matters because enterprise outcomes depend on more than a successful tool installation. The team must make provide reproducible workspaces with quotas, images, data access and lifecycle controls. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of kubeflow notebooks and workspace isolation.
- Design the capability using Kubeflow Notebooks, PVCs, namespaces with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build a governed data-science workspace.

Core Concepts

Provide reproducible workspaces with quotas and problem framing

Provide reproducible workspaces with quotas and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For kubeflow notebooks and workspace isolation, record assumptions and the evidence that would disprove them.

images architecture and operating boundaries

The architecture for images architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

data access and lifecycle controls. automation, policy and evidence

Automation for data access and lifecycle controls. automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Kubeflow Notebooks and Workspace Isolation failure handling and continuous improvement

Continuous improvement for kubeflow notebooks and workspace isolation failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubeflow Notebooks	Implement the core capability and define its service contract.
Automation and integration	PVCs	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	namespaces	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubeflow Notebooks	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	PVCs	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	namespaces	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: kubeflow_notebooks_and_workspace
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 22

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 22: Build a governed data-science workspace

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for kubeflow notebooks and workspace isolation.
2. Create the smallest safe implementation or detailed simulation using Kubeflow Notebooks, PVCs, namespaces.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Kubeflow Notebooks and Workspace Isolation?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Kubeflow Notebooks and Workspace Isolation is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 23: Kubeflow Pipelines Engineering

Define portable DAGs, components, caching, artifacts and metadata.

Why this chapter matters: Kubeflow Pipelines Engineering matters because enterprise outcomes depend on more than a successful tool installation. The team must make define portable dags, components, caching, artifacts and metadata. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of kubeflow pipelines engineering.
- Design the capability using Kubeflow Pipelines SDK, components with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build and compile a reusable pipeline.

Core Concepts

Define portable DAGs and problem framing

Define portable DAGs and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For kubeflow pipelines engineering, record assumptions and the evidence that would disprove them.

components architecture and operating boundaries

The architecture for components architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

caching automation, policy and evidence

Automation for caching automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

artifacts and metadata. failure handling and continuous improvement

Continuous improvement for artifacts and metadata. failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubeflow Pipelines SDK	Implement the core capability and define its service contract.
Automation and integration	components	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Kubeflow Pipelines SDK	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	components	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Kubeflow Pipelines SDK	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	components	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: kubeflow_pipelines_engineering
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 23

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 23: Build and compile a reusable pipeline

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for kubeflow pipelines engineering.
2. Create the smallest safe implementation or detailed simulation using Kubeflow Pipelines SDK, components.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Kubeflow Pipelines Engineering?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Kubeflow Pipelines Engineering is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 24: Kubeflow Trainer and Distributed Jobs

Run fault-aware distributed training with scalable worker topologies.

Why this chapter matters: Kubeflow Trainer and Distributed Jobs matters because enterprise outcomes depend on more than a successful tool installation. The team must make run fault-aware distributed training with scalable worker topologies. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of kubeflow trainer and distributed jobs.
- Design the capability using Kubeflow Trainer, PyTorch, MPI with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: launch and recover a distributed training job.

Core Concepts

Run fault-aware distributed training with scalable worker topologies. and problem framing

Run fault-aware distributed training with scalable worker topologies. and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For kubeflow trainer and distributed jobs, record assumptions and the evidence that would disprove them.

Kubeflow Trainer and Distributed Jobs architecture and operating boundaries

The architecture for kubeflow trainer and distributed jobs architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Kubeflow Trainer and Distributed Jobs automation, policy and evidence

Automation for kubeflow trainer and distributed jobs automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Kubeflow Trainer and Distributed Jobs failure handling and continuous improvement

Continuous improvement for kubeflow trainer and distributed jobs failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubeflow Trainer	Implement the core capability and define its service contract.
Automation and integration	PyTorch	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	MPI	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubeflow Trainer	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	PyTorch	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	MPI	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: kubeflow_trainer_and_distributed
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 24

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 24: Launch and recover a distributed training job

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for kubeflow trainer and distributed jobs.
2. Create the smallest safe implementation or detailed simulation using Kubeflow Trainer, PyTorch, MPI.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Kubeflow Trainer and Distributed Jobs?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Kubeflow Trainer and Distributed Jobs is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 25: Katib Hyperparameter Optimization

Automate search spaces, objectives, early stopping and experiment governance.

Why this chapter matters: Katib Hyperparameter Optimization matters because enterprise outcomes depend on more than a successful tool installation. The team must make automate search spaces, objectives, early stopping and experiment governance. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of katib hyperparameter optimization.
- Design the capability using Kubeflow Katib, trial templates with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: run a bounded hyperparameter experiment.

Core Concepts

Automate search spaces and problem framing

Automate search spaces and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For katib hyperparameter optimization, record assumptions and the evidence that would disprove them.

objectives architecture and operating boundaries

The architecture for objectives architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

early stopping and experiment governance. automation, policy and evidence

Automation for early stopping and experiment governance. automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Katib Hyperparameter Optimization failure handling and continuous improvement

Continuous improvement for katib hyperparameter optimization failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubeflow Katib	Implement the core capability and define its service contract.
Automation and integration	trial templates	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Kubeflow Katib	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	trial templates	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Kubeflow Katib	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	trial templates	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```
capability: katib_hyperparameter_optimizatio
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 25
```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 25: Run a bounded hyperparameter experiment

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for katib hyperparameter optimization.
2. Create the smallest safe implementation or detailed simulation using Kubeflow Katib, trial templates.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Katib Hyperparameter Optimization?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Katib Hyperparameter Optimization is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 26: GPU Scheduling, Quotas and Topology

Allocate accelerators fairly while respecting locality, fragmentation and priorities.

Why this chapter matters: GPU Scheduling, Quotas and Topology matters because enterprise outcomes depend on more than a successful tool installation. The team must make allocate accelerators fairly while respecting locality, fragmentation and priorities. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of gpu scheduling, quotas and topology.
- Design the capability using Kubernetes device plugins, Kueue, Volcano with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: design a gpu quota and scheduling policy.

Core Concepts

Allocate accelerators fairly while respecting locality and problem framing

Allocate accelerators fairly while respecting locality and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For gpu scheduling, quotas and topology, record assumptions and the evidence that would disprove them.

fragmentation and priorities. architecture and operating boundaries

The architecture for fragmentation and priorities. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

GPU Scheduling, Quotas and Topology automation, policy and evidence

Automation for gpu scheduling, quotas and topology automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

GPU Scheduling, Quotas and Topology failure handling and continuous improvement

Continuous improvement for gpu scheduling, quotas and topology failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubernetes device plugins	Implement the core capability and define its service contract.
Automation and integration	Kueue	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Volcano	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubernetes device plugins	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	Kueue	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	Volcano	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: gpu_scheduling_quotas_and_topolo
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 26

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 26: Design a GPU quota and scheduling policy

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for gpu scheduling, quotas and topology.
2. Create the smallest safe implementation or detailed simulation using Kubernetes device plugins, Kueue, Volcano.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design GPU Scheduling, Quotas and Topology?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- GPU Scheduling, Quotas and Topology is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 27: Ray Data for AI Workloads

Scale preprocessing, batch inference and streaming data loading across CPU and GPU.

Why this chapter matters: Ray Data for AI Workloads matters because enterprise outcomes depend on more than a successful tool installation. The team must make scale preprocessing, batch inference and streaming data loading across cpu and gpu. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of ray data for ai workloads.
- Design the capability using Ray Data, object storage, Parquet with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: build a scalable preprocessing pipeline.

Core Concepts

Scale preprocessing and problem framing

Scale preprocessing and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For ray data for ai workloads, record assumptions and the evidence that would disprove them.

batch inference and streaming data loading across CPU and GPU. architecture and operating boundaries

The architecture for batch inference and streaming data loading across cpu and gpu. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Ray Data for AI Workloads automation, policy and evidence

Automation for ray data for ai workloads automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Ray Data for AI Workloads failure handling and continuous improvement

Continuous improvement for ray data for ai workloads failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Ray Data	Implement the core capability and define its service contract.
Automation and integration	object storage	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Parquet	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Ray Data	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	object storage	Define upgrade, rollback, backup, restore and deprecation behavior.

Capability	Tool or pattern	Enterprise use
Portfolio proof	Parquet	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: ray_data_for_ai_workloads
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 27

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 27: Build a scalable preprocessing pipeline

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for ray data for ai workloads.
2. Create the smallest safe implementation or detailed simulation using Ray Data, object storage, Parquet.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Ray Data for AI Workloads?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Ray Data for AI Workloads is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 28: Ray Train for Distributed Training

Scale training code with checkpointing, recovery and resource-aware workers.

Why this chapter matters: Ray Train for Distributed Training matters because enterprise outcomes depend on more than a successful tool installation. The team must make scale training code with checkpointing, recovery and resource-aware workers. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of ray train for distributed training.
- Design the capability using Ray Train, PyTorch/TensorFlow, checkpoints with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: run a distributed training and recovery test.

Core Concepts

Scale training code with checkpointing and problem framing

Scale training code with checkpointing and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For ray train for distributed training, record assumptions and the evidence that would disprove them.

recovery and resource-aware workers. architecture and operating boundaries

The architecture for recovery and resource-aware workers. architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

Ray Train for Distributed Training automation, policy and evidence

Automation for ray train for distributed training automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

Ray Train for Distributed Training failure handling and continuous improvement

Continuous improvement for ray train for distributed training failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Ray Train	Implement the core capability and define its service contract.
Automation and integration	PyTorch/TensorFlow	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	checkpoints	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Ray Train	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	PyTorch/TensorFlow	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	checkpoints	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: ray_train_for_distributed_traini
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 28

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 28: Run a distributed training and recovery test

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for ray train for distributed training.
2. Create the smallest safe implementation or detailed simulation using Ray Train, PyTorch/TensorFlow, checkpoints.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Ray Train for Distributed Training?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Ray Train for Distributed Training is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 29: KubeRay Operations and Cluster Lifecycle

Operate Ray clusters, jobs and services with autoscaling, monitoring and upgrades.

Why this chapter matters: KubeRay Operations and Cluster Lifecycle matters because enterprise outcomes depend on more than a successful tool installation. The team must make operate ray clusters, jobs and services with autoscaling, monitoring and upgrades. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of kuberay operations and cluster lifecycle.
- Design the capability using KubeRay, RayJob, RayService, Kubernetes with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: deploy and upgrade a kuberay workload.

Core Concepts

Operate Ray clusters and problem framing

Operate Ray clusters and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For kuberay operations and cluster lifecycle, record assumptions and the evidence that would disprove them.

jobs and services with autoscaling architecture and operating boundaries

The architecture for jobs and services with autoscaling architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

monitoring and upgrades. automation, policy and evidence

Automation for monitoring and upgrades. automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

KubeRay Operations and Cluster Lifecycle failure handling and continuous improvement

Continuous improvement for kuberay operations and cluster lifecycle failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	KubeRay	Implement the core capability and define its service contract.
Automation and integration	RayJob	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	RayService	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	Kubernetes	Expose health, usage, audit and user-outcome signals.

Capability	Tool or pattern	Enterprise use
Lifecycle and recovery	KubeRay	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	RayJob	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: kuberay_operations_and_cluster_l
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 29

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 29: Deploy and upgrade a KubeRay workload

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for kuberay operations and cluster lifecycle.
2. Create the smallest safe implementation or detailed simulation using KubeRay, RayJob, RayService, Kubernetes.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design KubeRay Operations and Cluster Lifecycle?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- KubeRay Operations and Cluster Lifecycle is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Chapter 30: Distributed AI Platform Capstone

Combine workspaces, pipelines, training, scheduling, data and operational controls.

Why this chapter matters: Distributed AI Platform Capstone matters because enterprise outcomes depend on more than a successful tool installation. The team must make combine workspaces, pipelines, training, scheduling, data and operational controls. explicit, assign ownership, automate repeatable controls, and retain evidence that the design behaves correctly during normal use, denied use and failure.

Learning Objectives

- Explain the purpose, consumers and enterprise trade-offs of distributed ai platform capstone.
- Design the capability using Kubeflow, Ray, Kueue, OpenTelemetry with explicit ownership and failure boundaries.
- Implement a version-controlled workflow with identity, policy, telemetry and lifecycle controls.
- Troubleshoot common failures and produce evidence suitable for operations, security and architecture review.
- Complete the lab: deliver a multi-tenant distributed ai platform.

Core Concepts

Combine workspaces and problem framing

Combine workspaces and problem framing starts with a clear user or business problem. Define the desired outcome, the actors, the service boundary and the constraints before selecting implementation technology. For distributed ai platform capstone, record assumptions and the evidence that would disprove them.

pipelines architecture and operating boundaries

The architecture for pipelines architecture and operating boundaries must separate control, execution, data and evidence paths. Document trust boundaries, dependencies, quotas, failure domains and who owns each operational decision. Abstraction should reduce repeated work without concealing critical behavior.

training automation, policy and evidence

Automation for training automation, policy and evidence must be versioned, reviewed and testable. Include schema validation, policy, dry runs or plans, staged promotion, post-change verification and rollback. Manual emergency actions require an owner, timestamp, reason and expiry.

scheduling failure handling and continuous improvement

Continuous improvement for scheduling failure handling and continuous improvement requires observable user outcomes and bounded failure exercises. Test one expected path, one denied path and one dependency failure; then update documentation, alerts and product backlog from the evidence.

Capability and Tool Map

Capability	Tool or pattern	Enterprise use
Primary platform or framework	Kubeflow	Implement the core capability and define its service contract.
Automation and integration	Ray	Connect the capability to version-controlled workflows and APIs.
Policy, security or governance	Kueue	Constrain behavior with least privilege, validation and exceptions.
Telemetry and operational evidence	OpenTelemetry	Expose health, usage, audit and user-outcome signals.
Lifecycle and recovery	Kubeflow	Define upgrade, rollback, backup, restore and deprecation behavior.
Portfolio proof	Ray	Produce sanitized architecture, test and decision evidence.

Architecture and Operating Decisions

- Start with an explicit consumer outcome and service contract; do not expose a tool merely because it is available.
- Separate control-plane authority from workload execution and minimize the maximum blast radius of credentials, state and change.
- Prefer short-lived identity, immutable artifacts, versioned configuration and automated policy over persistent credentials and console-only workflows.
- Define SLOs, ownership, support, cost allocation, recovery, upgrade and deprecation before broad adoption.
- Preserve an escape hatch with accountable review when the golden path cannot meet a valid requirement.

Reference Service Contract

```

capability: distributed_ai_platform_capstone
owner: platform-or-product-team
consumers: [developer, operator, security, finance]
contract:
  inputs: [request, metadata, policy_context]
  outputs: [resource_or_service, status, evidence]
controls:
  identity: short_lived_and_least_privilege
  policy: versioned_and_tested
  telemetry: metrics_logs_traces_audit
  lifecycle: provision_upgrade_recover_deprecate
validation:
  - approved_path
  - denied_path
  - dependency_failure
  - rollback_or_restore
  - cost_and_ownership_check
volume: 14C
chapter: 30

```

Production Readiness Checklist

- [] Consumers, owner, support tier, lifecycle state and service-level objectives are recorded.
- [] Identity, tenant boundaries, data classification, network paths and secrets are tested.
- [] Configuration and policy are versioned, reviewed, promoted and observable.
- [] Failure, rollback or restore, capacity, quota and dependency behavior have been exercised.
- [] Usage, adoption, user outcomes, cost allocation and exception expiry are measurable.

Troubleshooting Matrix

Symptom	Likely cause	Corrective action	Validation evidence
Request remains pending	Controller, runner or dependency cannot reconcile the desired state	Inspect status conditions, events, queue depth and dependency health; identify the first failing boundary	Successful reconciliation plus latency and audit evidence
User is denied unexpectedly	Identity mapping, policy context or ownership metadata is incomplete	Evaluate effective identity and policy inputs; correct the narrowest cause and rerun denied/allowed tests	Expected allow and deny results with access logs
Change succeeds but user outcome degrades	Technical success criteria do not represent service behavior	Compare SLO, dependency and user-journey signals before and after the change; roll back if needed	Restored SLO and documented causal evidence
Cost or capacity grows unexpectedly	Unbounded defaults, leaks, retries, idle resources or poor allocation	Break down consumption by tenant and unit; apply quota, TTL, scaling or design changes	Reduced unit cost with reliability and quality unchanged

Hands-on Lab 30: Deliver a multi-tenant distributed AI platform

Prerequisites

- Disposable local or cloud sandbox with budgets and expiry controls.
- Git, containers, Kubernetes or a documented simulation as appropriate.
- Synthetic identities and data only; do not use employer or customer secrets.

Implementation

1. Define fictional requirements, users, ownership and acceptance criteria for distributed ai platform capstone.
2. Create the smallest safe implementation or detailed simulation using Kubeflow, Ray, Kueue, OpenTelemetry.
3. Add identity, policy, metadata, telemetry, quota, cost and lifecycle controls.
4. Run one successful user journey and one intentionally denied action.
5. Inject a bounded dependency or configuration failure, diagnose it, recover and record the timeline.
6. Publish a sanitized architecture, runbook, evidence pack and improvement backlog; then clean up all resources.

Validation Evidence

- Architecture and service contract committed to version control.
- Commands, test results, screenshots or exported telemetry for allow, deny and failure scenarios.
- Recovery, cost, ownership, security and cleanup evidence with timestamps.

Interview Questions and Model Answers

1. How would you design Distributed AI Platform Capstone?

Begin with consumers, outcomes, constraints and ownership. Define the service contract and failure boundaries, choose the minimum viable architecture, automate policy and lifecycle, then prove normal, denied and recovery behavior.

2. Which trade-offs would you document?

Document speed versus control, abstraction versus transparency, standardization versus flexibility, managed service versus portability, centralization versus team autonomy, and cost versus resilience.

3. What production evidence is required?

Require version-controlled architecture and configuration, effective-access results, policy tests, user-journey telemetry, SLOs, rollback or restore evidence, ownership, cost allocation and a clean teardown or deprecation path.

4. How would you improve adoption?

Use developer research, remove the largest friction, ship a narrow golden path, publish support and escape hatches, measure time-to-value and satisfaction, and iterate with representative product teams.

Chapter Summary

- Distributed AI Platform Capstone is an owned product or operating capability, not a one-time installation.
- Enterprise readiness requires explicit contracts, boundaries, policy, telemetry and lifecycle controls.
- Adoption and reliability improve when normal, denied and failure behavior are tested with users.
- The lab artifact should show reasoning and evidence, not confidential environment details.

Distributed AI Platform Reference

Use this reference to review product outcomes, architecture boundaries, operating controls and evidence required for production approval.

Domain	Required outcome	Validation evidence
Product and users	Named consumers, outcome hypothesis, roadmap and lifecycle owner	Research summary, service contract and adoption data
Architecture	Control, execution, data and evidence boundaries with bounded blast radius	Diagrams, ADRs, dependency and failure tests
Identity and security	Least privilege, keyless workloads, protected administration and auditable exceptions	Effective-access, denied-action and audit evidence
Delivery and policy	Versioned automation, policy, promotion, rollback and provenance	Pipeline, policy and artifact verification results
Reliability and operations	SLOs, telemetry, capacity, incident response, recovery and deprecation	Dashboards, game-day timeline and restore evidence
Data and AI governance	Ownership, classification, quality, privacy, lineage and responsible use	Catalog, tests, risk and approval records
Economics and value	Allocated cost, unit economics, investment rationale and optimization safety	Cost model, outcome metrics and decision log
People and governance	Clear decisions, support, skills, change plan and feedback loops	RACI, operating cadence, survey and action backlog

Completion Checklist

- [] I can explain every chapter in Volume 14C using architecture, operations and business language.
- [] I can implement or simulate at least one complete service with identity, policy, telemetry and lifecycle controls.
- [] I can demonstrate expected, denied and failure behavior with time-stamped evidence.
- [] I can discuss security, reliability, cost, adoption and organizational trade-offs.
- [] I can produce a sanitized portfolio artifact and defend the design in an interview or review board.

Additional Portfolio Projects

1. Service catalog and ownership quality automation.
2. Keyless CI/CD and workload-identity reference implementation.
3. Policy promotion pipeline with expiring exceptions.
4. Golden path or pipeline template with user-journey telemetry.
5. SLO dashboard and bounded failure game day.
6. Cost allocation and unit-economics review.
7. Security threat model and evidence-based remediation plan.
8. Backup, restore, failover and failback exercise.
9. Migration or adoption factory for three fictional teams.
10. Executive architecture or investment memo with measured outcomes.

Practice Plan

Days	Focus	Required output
1-5	Concepts, users and current-state discovery	Research notes, value stream and capability map
6-10	Architecture and service contracts	Diagrams, API/schema and ADRs
11-16	Implementation and automation	Working sandbox or detailed simulation
17-21	Security, policy and governance	Threat model, policy tests and exceptions
22-26	Reliability, telemetry and recovery	SLOs, dashboards and failure drill
27-31	Cost, capacity and lifecycle	Unit cost, quota and deprecation plan
32-36	Adoption, people and operating model	RACI, support and change plan
37-41	Troubleshooting and interview practice	Timed investigation and interview notes
42-45	Capstone and portfolio	Final defense, cleanup and sanitized evidence pack

Glossary

Term	Definition
Control plane	The APIs and controllers that accept desired state and coordinate execution.
Developer portal	A user interface that aggregates catalog, documentation, workflows and operational context.
Internal developer platform	The integrated products, APIs, automation and processes that support software delivery.
Golden path	An opinionated, supported default journey for a common engineering task.
Data product	An owned data asset or interface with consumers, quality, documentation and lifecycle.
MLOps	Practices and systems for reproducible, governed delivery and operation of machine-learning models.
LLMOps	Lifecycle, evaluation, deployment, observability, safety and cost practices for LLM applications.
Sociotechnical system	A system whose outcomes emerge from interactions among people, process, incentives and technology.
Unit economics	Cost and value measured per meaningful business or technical output.
Error budget	The allowed unreliability implied by a service-level objective.
Policy as code	Machine-testable policy stored, reviewed and promoted through version control.
Evidence pack	Sanitized architecture, configuration, test, telemetry and decision records used to prove readiness.

Official Documentation Index

The following primary or official sources were consulted for this 2026 edition. Verify current versions, service behavior, prices, limits and certification objectives before use.

- [MLflow Documentation](#)
- [MLflow Tracking](#)
- [MLflow Model Registry](#)
- [Kubeflow Introduction](#)

- [Kubeflow Pipelines](#)
- [Kubeflow Trainer](#)
- [KServe Documentation](#)
- [Ray Documentation](#)
- [Ray Serve](#)
- [NIST AI Risk Management Framework](#)

Volume 14C completion: preserve your architecture, code or simulation, test evidence, recovery timeline, cost/value review and sanitized final report as one portfolio package.